

ANOOP RAJENDRA

anoop.rajendra@gmail.com · (619) 630 5264

GitHub · LinkedIn · twitter

Summary

I design, architect, and build large scale Linux clusters. I currently manage a group that automates installation and deployment of Linux on clusters of hundreds of compute and storage servers. I've worked in every part of the Linux OS dealing with network, storage, and compute stacks. I'd like to work in a position where I can utilize my knowledge, expertise, and experience in the field of HPC, Clusters, Cloud, and Storage systems, and apply it, along with new and cutting edge technologies, to solve extremely large scale distributed computational and storage problems.

Skills

Technologies

Cluster Computing - Stacki, Open Grid Scheduler · Cloud Computing - AWS, GCP
Operating Systems - Linux, Solaris, *nix · Databases - MySQL, Redis
OS Distributions - RedHat, SuSE, CentOS · Web Services - Apache, Django

Development

Languages - Python, C, BASH · Git · Make · Jenkins

Experience

Teradata

Senior Engineering Manager · 07/2017 - Present · San Diego

Joined Teradata as part of Teradata's acquisition of StackIQ

Responsibilities

- Manage a 10 person Scrum Team for Infrastructure Automation
- Manage core Stacki framework development
- Manage SNAPI development - orchestration engine that calls Stacki
- Work with CI/CD team to build pipelines for software build/test/release/delivery lifecycle
- Liaising with external teams to help them adapt their software projects to Stacki
- Support external teams with Stacki deployments

Accomplishments

- Drove adoption of Stacki across the company
- Created standardized bare-metal provisioning process for Vantage - Teradata's flagship product
- Created Vantage Deployment Team to support bare-metal deployments for internal customers
- Drove standardized packaging for legacy database packages
- Won Product & Technology Award for Software Delivery in 2018

StackIQ

Director, Enterprise Software Development · 11/2015 - 06/2017 · San Diego, CA

Senior Software Developer · 10/2011 - 10/2015 · San Diego, CA

- Core developer on Stacki - Open Source Cluster management tool
- Core developer on Stacki Enterprise (the commercial version of Stacki)
- Provide installation and support services, software customization, and support for new, and existing customers
- Pallet development - additional software modules of generic third-party software for Stacki and Stacki Enterprise
- Project management for custom pallet development
- Developed pallets to support Big Data apps such as Hortonworks, MapR, Cloudera
- Developed pallets to support CloudStack, VMware ESXi, OpenStack
- Developed automation of Storage Controller configuration

San Diego Supercomputer Center

Programmer/Analyst III · 11/2005 - 10/2011 · San Diego, CA

- Core Developer on NPACI Rocks
- Ported Rocks code to the Solaris Operating System
- Used the Rocks infrastructure to install and maintain compute clusters and storage clusters
- Designed, developed, and extended the Rocks core platform, and additional modules called Rolls
- Adapted Bio-Informatics applications to the Rocks Platform

Fermi National Accelerator Laboratory

Graduate Research Assistant · 09/2004 - 09/2005 · Batavia, IL

- Worked on a Grid Computing Project called the SAM-Grid Project
- Core developer on the JIM (Job & Information Management) system

Education

MS, Computer Science and Engineering

UT Arlington · Arlington, TX · 08/2003 - 12/2005

Area of interest - Clusters and Grid computing

Thesis - "Integration of the SAM-Grid Infrastructure to the D-Zero Data Reprocessing Effort"

GPA - 3.7

BE, Computer Science and Engineering

BMS College of Engineering · Bangalore, India · 10/1999 - 07/2003

Graduated with **First Class with Distinction**